

10 Blog Posts Every **AI Professional** Should Read

Most AI aspirants read too many tutorials and too few deep technical articles. A strong reading list should teach you to choose architectures, evaluate systems, and build real products.



10 Deep-Dives



7 Core Skills



All Free To Read

The real value — these essays don't just explain models. They teach judgment: how to choose, evaluate, debug, and ship.

Tutorials Teach Syntax. Essays Teach Judgment.

A tutorial shows you how to call an API. An essay from someone who shipped the system shows you why it works — and where it breaks.



The Habit

Bookmarking another "top 10 tools" listicle. It feels current — but it rarely changes how you think.



The Fix

Read the practitioners who wrote about production systems, evals, and failure modes — then apply one idea to your own project.

Why it matters — the next 10 slides are essays worth slowing down for. What separates you is whether you apply them.

The Illustrated Transformer

The clearest visual walkthrough of attention and transformer internals — still the best starting point years later.

→ jalammar.github.io/illustrated-transformer · Jay Alammar

 FOUNDATIONS



Best For

Understanding attention visually



Why It's Here

Turns a dense architecture into an intuitive picture

Don't just read it — take notes as you go, then explain the idea back in your own words.

LLM Powered Autonomous Agents

A foundational breakdown of planning, memory, and tool use — the building blocks every agent system still relies on.

→ lilianweng.github.io/posts/2023-06-23-agent · Lilian Weng

 AGENT DESIGN



Best For

Designing agent architectures



Why It's Here

The reference post most agent frameworks trace back to

Don't just read it — take notes as you go, then explain the idea back in your own words.

AI Engineering

How AI systems actually become products — with data, evaluation, and deployment treated as first-class problems.

→ huyenchip.com/2025/01/07/ai-engineering.html · Chip Huyen

 AI PRODUCT ENGINEERING



Best For

Moving from model to product



Why It's Here

Product-first thinking, not just model-first

Don't just read it — take notes as you go, then explain the idea back in your own words.

Your AI Product Needs Evals

The case for building evaluation and feedback loops before you touch the model again — most teams skip this and pay for it later.

→ hamel.dev/blog/posts/evals · Hamel Husain

 EVALUATION



Best For

Teams shipping without evals



Why It's Here

Why eval-first beats prompt-first

Don't just read it — take notes as you go, then explain the idea back in your own words.

Building Effective AI Agents

A clear-eyed comparison of workflows vs. autonomous agents — and why simplicity should always be the default.

→ anthropic.com/research/building-effective-agents · Anthropic

 AGENT DESIGN



Best For

Choosing workflow vs. agent



Why It's Here

Reusable decision framework, not hype

Don't just read it — take notes as you go, then explain the idea back in your own words.

Contextual Retrieval

A practical technique for improving retrieval by contextualizing chunks before embedding — fewer missing-context failures in RAG.

→ anthropic.com/news/contextual-retrieval · Anthropic

 RETRIEVAL & CONTEXT



Best For

RAG systems losing context



Why It's Here

A concrete fix, not just a concept

Don't just read it — take notes as you go, then explain the idea back in your own words.

Patterns for Building LLM Applications

A catalog of reusable patterns that connect raw model capability to actual product requirements.

→ eugeneyan.com/writing/llm-patterns · Eugene Yan

 REUSABLE PATTERNS



Best For

Designing LLM app architecture



Why It's Here

A pattern library you'll return to repeatedly

Don't just read it — take notes as you go, then explain the idea back in your own words.

Better Experiments With LLM Evals

Why offline evals alone aren't enough — and how to validate quality against real user outcomes with live experiments.

→ engineering.atspotify.com/2026/5/better-experiments-with-llm-evals · Spotify

 EVALUATION



Best For

Evals disconnected from real users



Why It's Here

Treats evals as a funnel, not a fork

Don't just read it — take notes as you go, then explain the idea back in your own words.

The Lethal Trifecta For AI Agents

A sharp framing of the exact combination — private data, untrusted content, external actions — that makes agents dangerous.

→ simonwillison.net/2025/Jun/16/the-lethal-trifecta · Simon Willison

 AI SAFETY



Best For

Connecting agents to real actions



Why It's Here

Names the risk pattern precisely, before you ship it

Don't just read it — take notes as you go, then explain the idea back in your own words.

Context Engineering

A grounded look at what actually belongs inside an LLM's context window — and how retrieval and state improve reliability.

→ louisbouchard.ai/learn-ai-engineering · Louis-François Bouchard

 **RETRIEVAL & CONTEXT**



Best For

Reliability issues from bad context



Why It's Here

Context as a design discipline, not an afterthought

Don't just read it — take notes as you go, then explain the idea back in your own words.

Seven Skill Areas, One Reading List

Every post above maps to a distinct part of the AI engineering stack.



Model Foundations

Architecture & attention intuition



Agent Design

Planning, memory & tool use



AI Product Engineering

Shipping models as real products



Evaluation

Offline evals & live experiments



Retrieval & Context

What belongs in the context window



AI Safety

Risk patterns before you ship

Plus — reusable LLM application patterns that tie every other skill area together.

Read With Intent, **Not Volume**

Ten essays won't help if you skim all ten. Pick one tied to a real problem you're facing right now.

1 Pick one, tied to a real problem

A shaky RAG pipeline, an agent that misbehaves, an eval gap — let the problem choose the post.

2 Take notes as you read

Not after. Capture the one idea that changes how you'd design something.

3 Explain it in your own words

If you can't restate it simply, you skimmed it — go back.

4 Apply it within a week

That's the point where a bookmark becomes an actual engineering habit.

Don't Just Read. **Apply.**

Pair this reading list with structured interview prep across every AI Engineering fundamental.

🔥 Master 550+ production-focused interview questions across 14 chapters — Python, RAG, AI Agents, LangGraph, Guardrails, AWS, System Design, LLMs, and more.

🌟 Complete AI interview guide: lnkd.in/g6KiByd8

👤 Follow Roy's AI Lab for more curated AI Engineering resources, every week.