
EXECUTIVE SERIES

Enterprise AI Economics

From Token Consumption to Trusted Business Outcomes

Why architecture, model portfolios, security, and compliance must be managed as one economic system

Prepared by

Prabhat Kumar

www.linkedin.com/in/eprabhat

Akriti Srivastava

<https://www.linkedin.com/in/akriti-srivastava-7857b6216/>

Date

September 2026

Version

1.0.2

Evidence / References

Public primary sources only (NIST, OWASP, European Commission, Government of India, Salesforce, Sierra).

Contents

AI economics is larger than token economics	2
The commercial shift from access to consumption and outcomes	3
Architecture is the economic engine	4
1. Context precision.....	4
2. Model routing.....	4
3. Memory and reuse	5
4. Agent and tool design	5
5. End-to-end cost observability	5
Model selection as portfolio management	6
The TRUST Framework for AI Security and Compliance	6
T: Track and classify.....	6
R: Restrict access and agency	6
U: Understand and minimize data exposure.....	7
S: Secure models, context, and execution	7
T: Test, document, and govern continuously	8
Regulatory compliance is use-case specific	8
A board-level AI scorecard.....	9
Economic measures	9
Value measures.....	9
Quality measures	10
Trust measures.....	10
References	11
Disclaimer	11

Enterprise AI is entering a more consequential phase. Organizations are progressing from standalone chat interfaces to systems that retrieve corporate knowledge, coordinate software tools, generate decisions, and execute multi-step business processes.

That transition changes the economics of AI. It also changes its risk profile.

A seemingly simple request can activate a chain of model calls, retrieved documents, system instructions, tool definitions, validation routines, stored memory, and generated outputs. Token prices are visible, but many of the architectural choices that determine total consumption remain hidden.

For executives, the relevant question is therefore no longer simply how much a model costs per million tokens. It is how efficiently, securely, and reliably the complete AI system converts model consumption into an accepted business outcome.

The central proposition of this article is:

Token price determines the rate paid for model consumption. Architecture determines the quantity consumed. Governance determines whether the activity is controlled and permissible. Business outcomes determine whether the expenditure created value.

AI economics is larger than token economics

Tokens are a legitimate unit of technical consumption. They represent portions of the information processed or generated by language models and give organizations a mechanism for tracking inference costs.

However, tokens do not measure whether a task was completed successfully.

The same volume of tokens might produce a high-quality research synthesis, resolve a customer request, generate an answer that requires substantial human correction, or be wasted on irrelevant context and repeated model calls.

Executives should therefore view AI economics as a value chain:

Infrastructure and tokens → model calls → completed tasks → accepted outcomes → measurable business value

Each stage introduces possible waste.




Figure 1. Original synthesis for this article, consistent with the value-chain framing in the text.

A system may retrieve more information than the model needs. It may send the same instructions with every request, repeatedly invoke tools, keep unnecessary history in the active context, or use a frontier model for a routine classification task. A cheaper model may also become expensive if it generates more retries, escalations, or human review.

The right financial objective is not the lowest token price. It is the lowest sustainable cost for an outcome that meets business, quality, security, and compliance requirements.

The commercial shift from access to consumption and outcomes

The software industry has moved through several pricing models. Traditional software was commonly priced through licenses or user seats. Cloud infrastructure popularized usage-based pricing. AI services have extended consumption pricing through tokens, requests, processing time, and agent interactions.

AI agents introduce another possibility: pricing linked to completed outcomes.

Sierra's published explanation of outcome-based pricing distinguishes it from both fixed-price software and consumption-based infrastructure. Sierra describes an approach in which charges are connected to specified results, including resolved support conversations, retained memberships, purchases, upsells, and cross-sells. Its article notes that, in most cases within its model, an unresolved conversation does not generate a charge. This is a vendor's description of its commercial approach, not an industry-wide definition.




Figure 2. Adapted from Sierra, "Outcome-based pricing for AI agents" (10 Dec 2024). Framing is the vendor's; not an industry-wide standard.

Bret Taylor, co-founder and CEO of Sierra and chair of the OpenAI board, has publicly discussed outcome-oriented software economics. Sierra's March 2026 program description says Taylor explained why he believes outcome-based pricing is the future of software business models and why he views a process, rather than a person, as the atomic unit of AI productivity. Because the public page summarizes a recorded discussion rather than presenting a complete quotation transcript in the retrieved text, these points are intentionally paraphrased here rather than presented as direct quotations.

Salesforce illustrates a related but evolving commercial model. Its September 2024 investor announcement presented Agentforce as a platform for agents that can analyze data, make decisions, and complete tasks. Salesforce's October 2024 investor announcement said Agentforce was generally available and described agents that could take action across business functions.

The lesson for buyers is not that outcome pricing will replace every other model. Consumption pricing aligns charges with usage, seat pricing provides predictability, and outcome pricing can align payment with a defined result. Enterprise agreements may combine committed capacity, usage limits, and performance-related components.

Publicly reported benchmarks give this model concrete shape. In customer-support use cases, several vendors now publish or disclose per-outcome rates: Intercom's Fin agent is priced at roughly \$0.99 per resolved conversation, Zendesk's AI agents at approximately \$1.50 to \$2.00 per automated resolution, and Salesforce's Agentforce Service Agent at about \$2.00 per conversation. Sierra negotiates a per-resolution rate

inside each enterprise contract rather than publishing one. A material design distinction is what counts as billable: some models charge only for a resolved outcome and treat escalations to a human as free, while others bill both resolutions and escalations, which weakens the alignment between price and value. Buyers should read each “resolution” definition carefully before signing.



Figure 3. Rates as publicly reported/disclosed (2024-2026): Intercom, Zendesk, Salesforce (via TechCrunch), and Sierra. “Resolution” definitions vary by vendor.

Before accepting an outcome-based contract, buyers must define four things precisely:

1. What constitutes a completed outcome?
2. Who verifies that the outcome is correct?
3. How are reversals, customer dissatisfaction, and human rework treated?
4. How are security incidents or non-compliant outcomes excluded?

Without those definitions, “pay for value” can become a commercial slogan rather than an auditable economic model.

Architecture is the economic engine

Model selection affects the price and performance of each call. Architecture determines how many calls occur, how much context each call receives, which tools are exposed, and how frequently work must be repeated.

1. Context precision

A model cannot independently determine whether every retrieved document is authoritative, current, or relevant. It processes the context supplied by the surrounding system.

An effective context architecture should aim to supply the smallest sufficient body of reliable evidence. This generally means applying access permissions, ranking sources, controlling duplication, identifying authoritative records, and limiting retrieval to the purpose of the task.

More context is not automatically better context. Excessive retrieval can increase cost while making it harder to distinguish relevant evidence from noise.

2. Model routing

Not every stage of an AI workflow requires the same degree of model capability.

Structured extraction, classification, formatting, and simple routing may be handled by smaller or specialized models. Complex synthesis, ambiguous reasoning, and higher-impact decisions may justify more capable models and stronger validation.

The appropriate principle is:

Select the least expensive model that consistently meets the defined quality, latency, security, and compliance threshold.

This is not the same as selecting the cheapest model. Model evaluations must measure the complete task, including retries, fallbacks, validation, human review, and correction.

3. Memory and reuse

An AI system that repeatedly reconstructs the same instructions, context, and reasoning plan pays the discovery cost again with every execution.

Approved workflow templates, cached results, validated knowledge objects, and reusable execution patterns can reduce that repetition. However, persistent AI memory can contain confidential data, personal information, inaccurate conclusions, and obsolete instructions.

Memory should therefore have an owner, purpose, classification, access policy, retention period, and correction or deletion mechanism.

4. Agent and tool design

Agents may be able to retrieve data, write to business systems, send communications, or trigger transactions. Their economic power and security risk grow together.

A well-controlled agent should receive only the tools and information necessary for its current task. Long-running workflows should manage state deliberately instead of adding every prior instruction and intermediate result to an ever-expanding context.

Material actions may require deterministic checks or human approval. Iteration, time, token, and spending limits can help prevent uncontrolled execution.

5. End-to-end cost observability

Token usage should be traced to a business process rather than measured only at the model endpoint.

A useful cost view includes:

- Model and embedding charges
- Retrieval and storage costs
- Tool and external API charges
- Validation and monitoring costs
- Human review and correction
- Failed or abandoned executions
- Security and compliance overhead
- Cost per completed and accepted result

This converts AI cost management from model procurement into workflow economics.

Model selection as portfolio management

An enterprise rarely needs a single model for every use case. It needs a governed portfolio whose components can be compared, routed, monitored, and replaced.

Evaluation should cover:

- Accuracy and task completion
- Cost per accepted outcome
- Latency and throughput
- Context-window requirements
- Data retention and provider-use terms
- Deployment location and residency
- Security and adversarial robustness
- Auditability and monitoring
- Portability and provider concentration
- Human-review requirements
- Migration and exit considerations

A model benchmark conducted in isolation cannot fully answer these questions. Models must be evaluated inside the intended workflow and against representative data, risk conditions, and acceptance criteria.

The TRUST Framework for AI Security and Compliance

The following **TRUST Framework** is an original executive synthesis for this article. It is not an official government, NIST, OWASP, or industry standard.

It uses public guidance from the NIST AI Risk Management Framework, the NIST Generative AI Profile, and the OWASP GenAI Security Project as reference points.

T: Track and classify

Maintain an inventory of AI systems and supporting components, including:

- Business and technical ownership
- Intended purpose and prohibited uses
- Models, providers, and versions
- Connected data and tools
- User and stakeholder groups
- Deployment locations
- Relevant jurisdictions
- Risk classification
- Current operational status

An organization cannot govern AI systems it cannot identify.

R: Restrict access and agency

Apply least privilege to users, agents, service identities, data sources, and tools.

Controls should be proportionate to the use case and may include permission-aware retrieval, separate workload identities, approved tool lists, action limits, environment separation, and human approval for consequential decisions.

The objective is to limit both the probability and potential impact of unauthorized or unintended actions.

U: Understand and minimize data exposure

AI data governance must cover more than model-training data. It should include prompts, uploaded content, retrieved passages, embeddings, generated outputs, feedback, logs, execution traces, and persistent memory.

Organizations should identify the purpose for which each category is processed, minimize unnecessary data, establish retention and deletion controls, and understand whether providers retain or reuse submitted information.

India's Press Information Bureau reported that the Digital Personal Data Protection Rules, 2025 were notified on 14 November 2025 (the underlying Ministry of Electronics and Information Technology Gazette notification, G.S.R. 846(E), is dated 13 November 2025) and described them as operationalizing the Digital Personal Data Protection Act, 2023. The government's announcement identified principles including consent and transparency, purpose limitation, data minimization, accuracy, storage limitation, security safeguards, and accountability. It also described an 18-month phased compliance timeline.

Application of those requirements depends on the processing activity and applicable law. Organizations should obtain qualified legal advice for specific implementations.

S: Secure models, context, and execution

As of 3 August 2026, OWASP describes its GenAI LLM Top 10 2026 as the latest community-driven guide to critical security risks affecting applications powered by large language models. OWASP says the edition includes updated rankings, expanded threat coverage, practical explanations, attack scenarios, and mitigations.

A security program should translate applicable threats into testable controls. Depending on the system, these may address:

- Prompt and indirect prompt injection
- Unauthorized disclosure of sensitive information
- Poisoned or untrusted data sources
- Unsafe handling of generated outputs
- Excessive agent authority
- Vulnerable dependencies and model supply chains
- Uncontrolled resource consumption
- Inadequate logging and incident response

OWASP guidance helps identify security risks, but using an OWASP checklist does not itself demonstrate legal or regulatory compliance.

The 2026 edition's ranking reflects several shifts relevant to the architecture principles above. Prompt Injection (LLM01) retains the top position, while Excessive Agency rose from sixth to third place—the largest upward move in the list—driven by production incidents in which agent outputs autonomously execute commands, invoke external APIs, or manage transactions. The former "System Prompt Leakage" category was retired and replaced by the broader "Hidden Context Exposure," recognizing that retrieved documents, agent memory,

tool responses, and application state carry as much confidentiality risk as the system prompt itself. These changes reinforce this article’s emphasis on context precision and constraining agent authority.

In September 2026, following publication of the 2026 Top 10 in August, the OWASP GenAI Security Project announced a complementary Agent Control Standard (ACS). The standard provides runtime policy-enforcement mechanisms and an Agent Bill of Materials (AgBOM) for inspecting, tracing, and instrumenting autonomous agents in production. For enterprises implementing the agent and tool-design controls described earlier, the ACS offers a machine-readable way to govern what agents are actually permitted to do at runtime.

OWASP GenAI LLM Top 10 — 2026 Edition

Prompt Injection holds #1; Excessive Agency is the largest riser; Hidden Context Exposure replaces System Prompt Leakage

LLM01	Prompt Injection	●
LLM02	Sensitive Information Disclosure	▲
LLM03	Excessive Agency	▲
LLM04	Supply Chain	▼
LLM05	Data and Model Poisoning	▼
LLM06	Unbounded Consumption	▲
LLM07	Misinformation	▲
LLM08	Hidden Context Exposure	✦
LLM09	Vector and Embedding Weaknesses	●
LLM10	Improper Output Handling	▼
▲ rose		▼ fell
		● held
		✦ new

Figure 5. OWASP GenAI Security Project, "OWASP GenAI LLM Top 10 2026" (published 3–4 Aug 2026). Movement shown vs. the 2025 edition.

T: Test, document, and govern continuously

NIST describes AI RMF 1.0 as a voluntary framework designed to help organizations incorporate trustworthiness considerations into the design, development, use, and evaluation of AI products, services, and systems. The framework organizes its core around four functions: Govern, Map, Measure, and Manage.

NIST’s Generative AI Profile, published in July 2024 as NIST AI 600-1, is a cross-sector companion to AI RMF 1.0. NIST says it is intended to help organizations identify risks specific to generative AI and select risk-management actions aligned with their priorities.

An evidence program can include intended-use documentation, evaluation results, security assessments, model and system records, human-oversight decisions, version histories, incident records, monitoring results, and approved exceptions.

Framework alignment should be treated as a risk-management discipline, not as automatic evidence of compliance with every applicable law.

Regulatory compliance is use-case specific

The EU AI Act applies progressively rather than through one universal commencement date.

The European Commission’s AI Act Service Desk states that the Act entered into force on 1 August 2024. General provisions, definitions, AI-literacy requirements, and prohibitions began applying on 2 February 2025.

Governance rules and obligations for general-purpose AI models began applying on 2 August 2025. The Service Desk identifies additional application dates extending through 2 August 2028.

Readers should note that this timeline was subsequently amended. The Digital Omnibus on AI (Regulation (EU) 2026/1744), which entered into force on 27 July 2026, deferred parts of the high-risk regime: stand-alone high-risk systems under Annex III now apply from 2 December 2027, and product-embedded high-risk systems under Annex I now apply from 2 August 2028. Notably, 2 August 2026 remained the general application date, and enforcement of the general-purpose AI, prohibition, transparency, and AI-literacy provisions was not deferred. Organizations should confirm the current Service Desk timeline for their specific obligations.

The Commission also provides an AI Act Explorer, a compliance checker, and access to the AI Act Service Desk. These public tools can assist stakeholders in understanding potential obligations, but they do not replace legal advice or a use-case-specific assessment.

The practical implication is that organizations should not label an entire platform “AI Act compliant” without identifying the system, its intended use, the organization’s legal role, and the provisions that apply.

EU AI Act — Phased Implementation Timeline

The Act applies progressively; high-risk deadlines were extended by Regulation (EU) 2026/1744



Figure 4. European Commission AI Act Service Desk implementation timeline, reflecting the Digital Omnibus on AI amendment (in force 27 Jul 2026).

A board-level AI scorecard

Boards and executive committees need a scorecard that connects economics with trust.

Economic measures

- Fully loaded cost per completed task
- Cost per accepted outcome
- Tokens and model calls per workflow
- Cost of validation and human review
- Budget variance
- Avoided cost through reuse

Value measures

- Cycle-time improvement
- Revenue or margin contribution
- Customer-resolution outcomes
- Capacity released for higher-value work

- Adoption and repeat usage

Quality measures

- Task-completion rate
- Accuracy against accepted evidence
- Exception and escalation rates
- Human override rate
- Rework required
- Performance drift

Trust measures

- Unauthorized-access attempts
- Sensitive-data incidents
- Unapproved agent actions
- Adversarial-test results
- Open risk exceptions
- Audit-evidence completeness
- Incident detection and containment

A proposed executive metric is **Trusted Outcome Yield**:

Verified business value ÷ fully loaded AI cost, provided defined quality, security, privacy, and compliance thresholds are satisfied

Enterprise AI economics cannot be reduced to token prices.

Tokens measure model consumption. Architecture determines how efficiently models, data, memory, and tools are assembled. Model portfolios balance capability, cost, and risk. Security constrains what systems and agents are permitted to access or execute. Compliance defines obligations that vary by jurisdiction and use case.

Business value is created only when the resulting outcome is useful, accepted, secure, and permissible.

The executive goal should therefore be neither maximum model capability nor minimum token consumption. It should be sustainable production of trusted business outcomes at an economically accountable cost.

References

1. Sierra, *Outcome-based pricing for AI agents*, 10 December 2024.
[Read Sierra's outcome-based pricing article \[sierra.ai\]](#)
2. Sierra, *Bret Taylor of Sierra on AI agents, outcome-based pricing, and the OpenAI board*, 10 March 2026.
[View Sierra's Bret Taylor program page \[sierra.ai\]](#)
3. Salesforce Investor Relations, *Salesforce Unveils Agentforce: What AI Was Meant to Be*, 12 September 2024.
[Read the Salesforce investor announcement \[investor.s...sforce.com\]](#)
4. Salesforce Investor Relations, *Salesforce's Agentforce Is Here: Trusted, Autonomous AI Agents to Scale Your Workforce*, 29 October 2024.
[Read the Agentforce general-availability announcement \[investor.s...sforce.com\]](#)
5. National Institute of Standards and Technology, *AI Risk Management Framework*.
[Access the official NIST AI RMF page \[nist.gov\]](#)
6. National Institute of Standards and Technology, *Artificial Intelligence Risk Management Framework, AI RMF 1.0*, NIST AI 100-1, January 2023.
[Read the official NIST AI RMF 1.0 publication \[nvlpubs.nist.gov\]](#)
7. National Institute of Standards and Technology, *Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile*, NIST AI 600-1, July 2024.
[View the official NIST publication page \[nist.gov\]](#)
8. National Institute of Standards and Technology, *NIST AI 600-1: Generative Artificial Intelligence Profile*.
[Read the official NIST Generative AI Profile \[nvlpubs.nist.gov\]](#)
9. OWASP GenAI Security Project, *OWASP GenAI LLM Top 10 2026*, 3 August 2026.
[Access the OWASP GenAI LLM Top 10 2026 \[genai.owasp.org\]](#)
10. OWASP GenAI Security Project, *Top 10 for LLM and GenAI Initiative*.
[View the OWASP LLM and GenAI initiative \[genai.owasp.org\]](#)
11. European Commission AI Act Service Desk, *Timeline for the Implementation of the EU AI Act*.
[View the official EU AI Act implementation timeline \[ai-act-ser....europa.eu\]](#)
12. European Commission AI Act Service Desk, *When does the AI Act go into effect?*
[Read the official implementation FAQ \[ai-act-ser....europa.eu\]](#)
13. European Commission, *AI Act Single Information Platform*.
[Access the European Commission AI Act platform \[ai-act-ser....europa.eu\]](#)
14. Government of India, Press Information Bureau, *Government notifies DPDP Rules to empower citizens and protect privacy*, 14 November 2025.
[Read the official Government of India announcement \[pib.gov.in\]](#)
15. Government of India, Press Information Bureau, *DPDP Rules, 2025 Notified: A Citizen-Centric Framework for Privacy Protection and Responsible Data Use*, 17 November 2025.
[Read the official PIB explainer \[pib.gov.in\]](#)

Disclaimer

Trusted Outcome Yield is an independently developed management construct introduced and defined in this article. Based on the author's review of the cited public sources, it is not an official metric of NIST, OWASP, the European Union, the Government of India, Sierra, or Salesforce. The article claims authorship of its original definition, wording, structure, and presentation, but does not claim exclusive copyright ownership of the underlying idea, method, or mathematical relationship. This statement does not constitute worldwide trademarks or intellectual-property clearance.